

# Actividad: Redes Neuronales Profundas

Tania Sayuri Guizado Hernández A01640092

*Importar TensorFlow*

```
In [2]: # Se cargan las librerías necesarias
import tensorflow as tf
import matplotlib.pyplot as plt
from tensorflow.keras import datasets, layers, models
```

*Descargamos el CIFAR-10 dataset y se prepara para su uso*

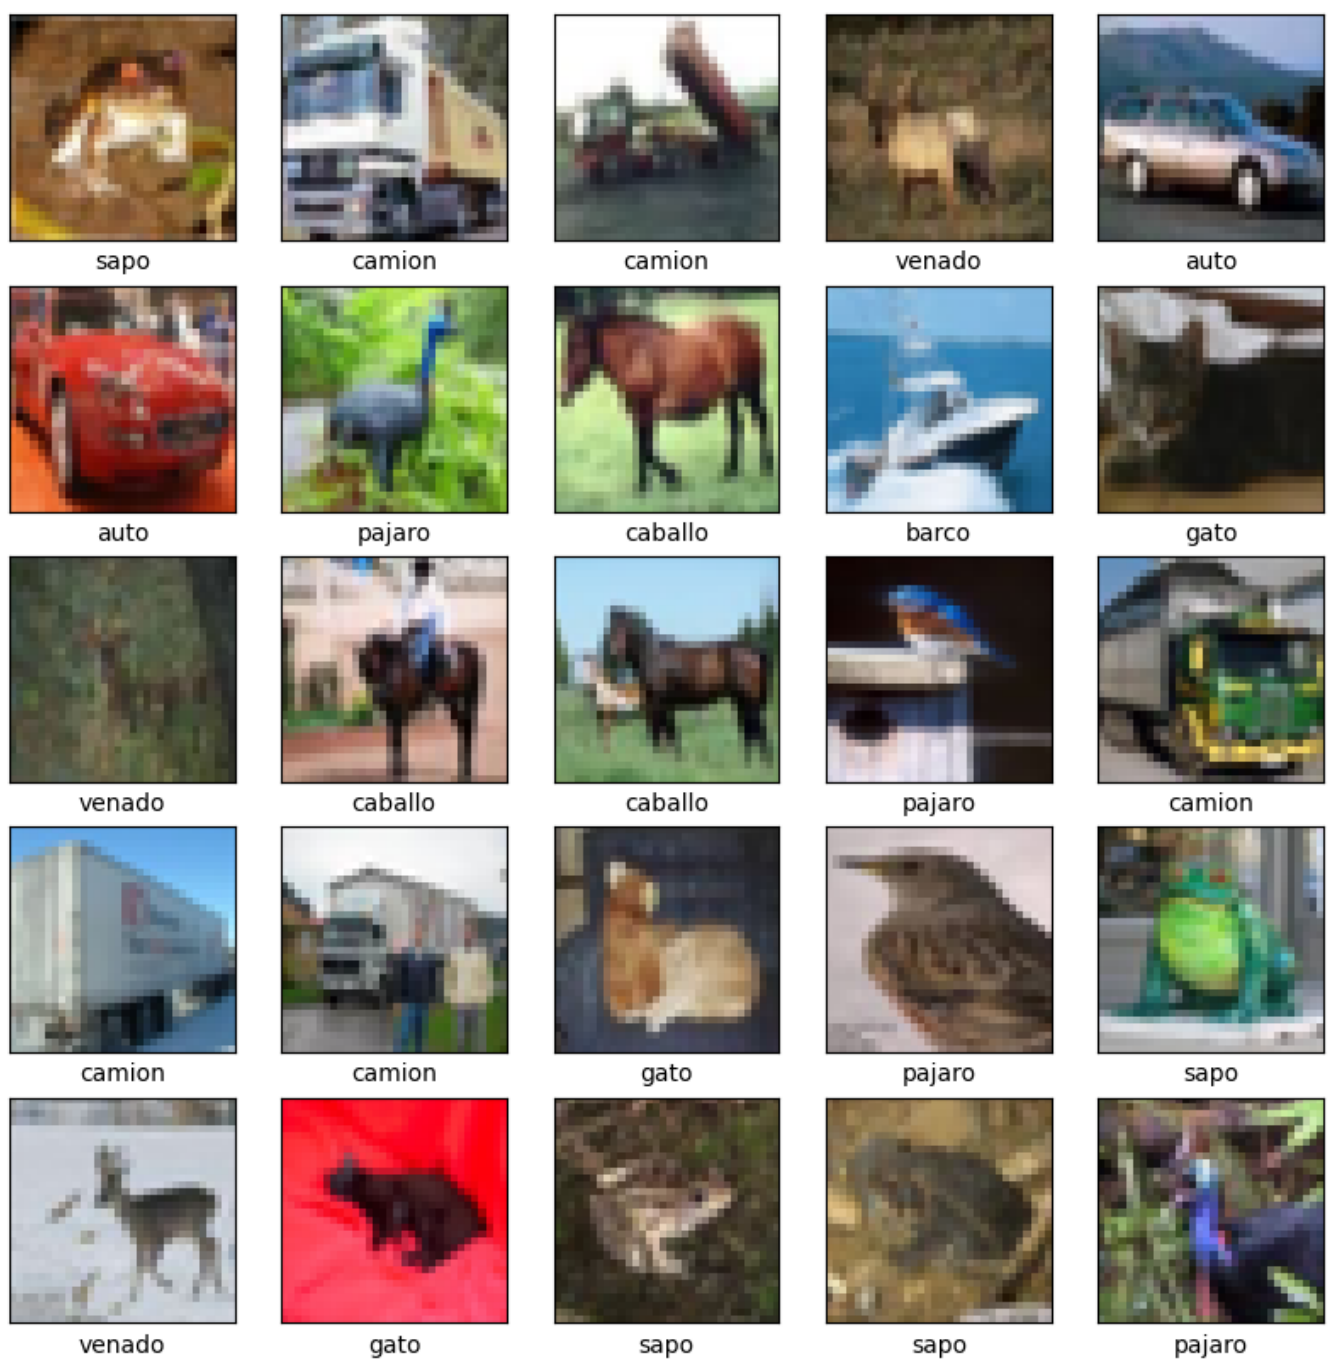
```
In [3]: # Se cargan los datos de CIFAR10
(train_images, train_labels), (test_images, test_labels) = datasets.cifar10.load_data()

# Normalizamos los valores de los píxeles entre 0 y 1
train_images, test_images = train_images/255.0, test_images/255.0
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>  
170498071/170498071 [=====] - 2s 0us/step

*Se crea una función para visualizar las imágenes*

```
In [5]: # Se guarda el nombre de las clases
class_names = ['avion', 'auto', 'pajaro', 'gato', 'venado', 'perro', 'sapo',
               'caballo', 'barco', 'camion']
plt.figure(figsize=(10,10))
for i in range(25):
    plt.subplot(5,5,i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(train_images[i])
    plt.xlabel(class_names[train_labels[i][0]])
plt.show()
```



### Capas de convolución

```
In [6]: # Se crea La red neuronal
model = models.Sequential()
model.add(layers.Conv2D(64, (3,3), activation='relu', input_shape=(32,32,3)))
model.add(layers.MaxPooling2D((2,2)))
model.add(layers.Conv2D(256, (3,3), activation='relu'))
model.add(layers.MaxPooling2D((2,2)))
model.add(layers.Conv2D(256, (3,3), activation='relu'))
model.add(layers.MaxPooling2D((2,2)))
```

### Arquitectura

```
In [7]: # Visualizamos La arquitectura de La red neuronal
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 30, 30, 64)	1792
max_pooling2d (MaxPooling2D)	(None, 15, 15, 64)	0
conv2d_1 (Conv2D)	(None, 13, 13, 256)	147712
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 256)	0
conv2d_2 (Conv2D)	(None, 4, 4, 256)	590080
max_pooling2d_2 (MaxPooling2D)	(None, 2, 2, 256)	0
=====		
Total params: 739584 (2.82 MB)		
Trainable params: 739584 (2.82 MB)		
Non-trainable params: 0 (0.00 Byte)		

### Capas densas

```
In [8]: # Se agregan las capas densas
model.add(layers.Flatten())
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dense(128, activation='sigmoid'))
```

```
In [9]: # Visualizamos la arquitectura del modelo
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 30, 30, 64)	1792
max_pooling2d (MaxPooling2D)	(None, 15, 15, 64)	0
conv2d_1 (Conv2D)	(None, 13, 13, 256)	147712
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 256)	0
conv2d_2 (Conv2D)	(None, 4, 4, 256)	590080
max_pooling2d_2 (MaxPooling2D)	(None, 2, 2, 256)	0
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 128)	131200
dense_1 (Dense)	(None, 128)	16512
=====		
Total params: 887296 (3.38 MB)		
Trainable params: 887296 (3.38 MB)		
Non-trainable params: 0 (0.00 Byte)		

### Compilación y entrenamiento

```
In [10]: model.compile(optimizer='adam',
                      loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
                      metrics=['accuracy'])
history = model.fit(train_images, train_labels, epochs=10, validation_data=(test_images, test_labels))
```

Epoch 1/10

/usr/local/lib/python3.10/dist-packages/keras/src/backend.py:5714: UserWarning: "`sparse\_categorical\_crossentropy` received `from\_logits=True`, but the `output` argument was produced by a Softmax activation and thus does not represent logits. Was this intended?

output, from\_logits = \_get\_logits(

1563/1563 [=====] - 317s 200ms/step - loss: 1.5088 - accuracy: 0.4578 - val\_loss: 1.1615 - val\_accuracy: 0.5846

Epoch 2/10

1563/1563 [=====] - 307s 197ms/step - loss: 1.0474 - accuracy: 0.6342 - val\_loss: 0.9670 - val\_accuracy: 0.6617

Epoch 3/10

1563/1563 [=====] - 302s 193ms/step - loss: 0.8600 - accuracy: 0.7029 - val\_loss: 0.8685 - val\_accuracy: 0.6952

Epoch 4/10

1563/1563 [=====] - 308s 197ms/step - loss: 0.7320 - accuracy: 0.7465 - val\_loss: 0.8203 - val\_accuracy: 0.7148

Epoch 5/10

1563/1563 [=====] - 305s 195ms/step - loss: 0.6194 - accuracy: 0.7849 - val\_loss: 0.8275 - val\_accuracy: 0.7262

Epoch 6/10

1563/1563 [=====] - 311s 199ms/step - loss: 0.5318 - accuracy: 0.8146 - val\_loss: 0.9106 - val\_accuracy: 0.7042

Epoch 7/10

1563/1563 [=====] - 307s 196ms/step - loss: 0.4547 - accuracy: 0.8416 - val\_loss: 0.8082 - val\_accuracy: 0.7386

Epoch 8/10

1563/1563 [=====] - 307s 197ms/step - loss: 0.3877 - accuracy: 0.8637 - val\_loss: 0.9435 - val\_accuracy: 0.7247

Epoch 9/10

1563/1563 [=====] - 309s 198ms/step - loss: 0.3298 - accuracy: 0.8848 - val\_loss: 0.8870 - val\_accuracy: 0.7381


Epoch 10/10

1563/1563 [=====] - 310s 199ms/step - loss: 0.2744 - accuracy: 0.9025 - val\_loss: 1.0280 - val\_accuracy: 0.7317

*Evaluation*

```
In [11]: plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label='val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.ylim([0.510, 1])
plt.legend(loc='lower right')
test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)
```

313/313 - 15s - loss: 1.0280 - accuracy: 0.7317 - 15s/epoch - 48ms/step



Se imprime el accuracy obteniendo un 0.7317 como resultado


```
In [12]: print(test_acc)
```

```
0.7317000031471252
```

Predicción

```
In [13]: n = 115 # Número de imagen
```

```
plt.figure(figsize=(2,2))
plt.imshow(test_images[n])
plt.xlabel(class_names[test_labels[n][0]])
plt.show()
```



```
In [14]: predictions = model.predict(test_images)
print(predictions[n])
```

```
import numpy as np
print('La imagen pertenece al grupo {} con una probabilidad de {:.2f}%'.format(
    class_names[np.argmax(predictions[n])], 100*np.max(predictions[n])))
```

313/313 [=====] - 16s 49ms/step  
[3.30988348e-01 1.16577685e-04 5.82631767e-01 9.69883442e-01  
9.99720991e-01 9.07494545e-01 4.40881878e-01 9.93510008e-01  
7.95677697e-05 9.70939845e-02 4.83510593e-11 6.93124308e-11  
1.51508302e-11 3.25139590e-12 9.65281024e-13 1.50641468e-12  
1.41729315e-12 1.74666520e-12 2.06333345e-12 3.12821609e-11  
5.11755189e-13 4.45048425e-12 2.99983482e-10 1.26959235e-10  
3.08201988e-12 5.44233278e-12 5.84768692e-12 2.09510686e-12  
1.17943308e-10 1.03099803e-11 1.96956302e-13 1.60939502e-12  
7.56235284e-14 5.15883891e-10 1.98884492e-10 1.24566189e-12  
1.69795449e-12 1.63234850e-13 5.50385995e-12 1.14981852e-12  
7.24131042e-13 1.16462439e-12 9.12997282e-12 4.56322317e-13  
1.31783924e-11 2.45881926e-11 1.09288559e-11 1.93267902e-11  
5.93407539e-13 8.38248637e-11 1.61588505e-12 5.33733578e-11  
2.80121629e-12 7.86221609e-12 1.48432551e-11 7.00218269e-12  
2.63367973e-11 3.77690657e-11 1.18020155e-12 9.39516585e-13  
3.56749461e-12 8.35964891e-12 4.67823891e-10 2.20528955e-12  
2.35063271e-12 1.71396189e-10 1.09321645e-12 1.83346556e-13  
3.19521102e-12 1.75570426e-11 1.93865150e-12 2.13940584e-12  
9.92696485e-13 2.28548387e-12 3.91090632e-12 2.52998837e-12  
1.29509884e-11 1.72147661e-11 6.66876346e-11 2.00748130e-14  
2.46465088e-11 7.45092703e-12 3.43555208e-13 1.43677110e-11  
7.85053272e-11 3.94431224e-11 5.82751417e-11 8.26109060e-12  
1.57250927e-12 4.66051586e-10 4.18428774e-13 2.49782318e-13  
1.18623462e-11 1.05715992e-11 1.40072502e-12 2.56888749e-13  
3.47688653e-12 1.18952443e-11 1.93775592e-13 2.98046332e-12  
1.39587245e-10 1.17309623e-12 3.08236076e-12 1.18831577e-11  
2.63732287e-12 1.67557569e-12 3.35080956e-11 4.29429305e-11  
7.90282510e-12 8.22577249e-13 9.61357036e-12 1.52765731e-10  
7.23971230e-12 2.87984653e-12 2.30806070e-13 4.73580593e-13  
6.22339040e-13 1.55178145e-11 1.43306658e-13 8.74147543e-12  
6.73509859e-13 5.66079301e-12 3.29799014e-12 3.70800873e-12  
3.79005800e-13 4.26858988e-13 6.33189638e-12 5.53454309e-13]

La imagen pertenece al grupo venado con una probabilidad de 99.97%